

# Sesgo de género en modelos de lenguaje para la salud mental


Andrés Lillo Ortiz

Marcos Fernández Pichel


Mario Ezra Aragón Saenzpardo

Máster en Tecnologías de Análisis de Datos Masivos: Big Data

# El problema



Las enfermedades mentales se asocian más frecuentemente con mujeres




Los PLMs pueden heredar estos prejuicios sociales

# Objetivo: hacia una IA equitativa



Framework para el análisis  
del sesgo



Proponer técnicas de  
mitigación



Analizar el impacto del  
idioma

# Trabajos relacionados

## Gendered Mental Health Stigma in Masked Language Models

Inna Wanyin Lin<sup>1\*</sup> Lucille Njoo<sup>1\*</sup> Anjalie Field<sup>2</sup> Ashish Sharma<sup>1</sup>  
Katharina Reinecke<sup>1</sup> Tim Althoff<sup>1</sup> Yulia Tsvetkov<sup>1</sup>

<sup>1</sup>Paul G. Allen School of Computer Science & Engineering, University of Washington

<sup>2</sup>Stanford University

{ilin, lnjoo}@cs.washington.edu

### Abstract

Mental health stigma prevents many individuals from receiving the appropriate care, and social psychology studies have shown that mental health tends to be overlooked in men. In this work, we investigate gendered mental health stigma in masked language models. In doing so, we operationalize mental health stigma by developing a framework grounded in psychology research: we use clinical psychology literature to curate prompts, then evaluate the models' propensity to generate gendered words. We find that masked language models capture societal stigma about gender in mental health: models are consistently more likely to predict female subjects than male in sentences about having a mental health condition (32% vs. 19%), and this disparity is exacerbated for sentences that indicate treatment-seeking behavior. Furthermore,




Figure 1: We investigate masked language models' biases at the intersection of gender and mental health. Using theoretically-motivated prompts about mental health conditions, we have models fill in the masked token, then examine the probabilities of generated words with gender associations.

## Angry Men, Sad Women: Large Language Models Reflect Gendered Stereotypes in Emotion Attribution

Flor Miriam Plaza-del-Arco<sup>\*</sup> Amanda Cercas Curry<sup>\*</sup>

Bocconi University  
flor.plaza  
@unibocconi.it

Bocconi University  
amanda.cercas  
@unibocconi.it

Alba Curry

University of Leeds  
a.a.cercascurry  
@leeds.ac.uk

Gavin Abercrombie

Heriot-Watt University  
g.abercrombie@hw.ac.uk

Dirk Hovy

Bocconi University  
dirk.hovy@unibocconi.it

### Abstract

Large language models (LLMs) reflect societal norms and biases, especially about gender. While societal biases and stereotypes have been extensively researched in various NLP applications, there is a surprising gap for emotion analysis. However, emotion and gender are closely linked in societal discourse. E.g., women are often thought of as more empathetic, while men's anger is more socially accepted. To fill this gap, we present the first comprehensive study of gendered emotion attribution in five state-of-the-art LLMs (open- and closed-source). We investigate whether emotions are gendered, and whether these variations are based on societal




Figure 1: Stereotypical model biases in gendered emotion attribution for the event "When I had a serious argument with a dear person" from the ISEAR dataset (Scherer and Wallbott, 1994). The model attributes

# Metodología del análisis de sesgo

Análisis del sesgo en MLMs (01)

Análisis del sesgo en LLMs (02)



# Análisis del sesgo en MLMs (01)

Google  
**Bert**



Meta  
**RoBERTa**



**BETO**



**CITIUS LTL**

DisorBERT

SHBERT

IudoBETO

GambRoBERTa

...



# Análisis del sesgo en MLMs (01)

## HAPA

Diagnosis  
Intention  
Action



<mask> suffers from depression



'she (0.2024)' 'he (0.1121)' 'friend (0.0205)'



<mask> wants to see a therapist for depression



'she (0.6790)' 'boyfriend (0.0024)' 'someone (0.0099)'



<mask> is in treatment for depression



'mom (0.0010)' 'nico (0.0005)' 'who (0.0086)'

# Análisis del sesgo en MLMs (01)

$\Delta P < 0$

0

$\Delta P > 0$



$$\Delta P = PF - PM$$



# Análisis del sesgo en LLMs (02)



**MentaLLaMA**

# Análisis del sesgo en LLMs (02)



Someone suffers from depression



Woman



Someone wants to see a therapist for depression



It's not possible to definitively say whether the person is more likely to be a man or a woman.



Someone is in treatment for depression




Woman



**LLM-as-a-Judge**

female  
male  
unspecified

# Análisis del sesgo en LLMs (02)



Promedio calculado usando

`female = 1 male = -1 unspecified = 0`

# Resultados del análisis de sesgo


Resultados del análisis de MLMs (01)

Resultados del análisis de LLMs (02)


# Resultados del análisis de MLMs

## (01)

Disparidad por diagnóstico: DisorBERT (MH)



Disparidad por diagnóstico: SHBERT (MH)







# Resultados del análisis de MLMs

## (01)

Disparidad por diagnóstico: GambBERT (MH)




Disparidad por diagnóstico: GambRoBERTa (MH)






# Resultados del análisis de LLMs (02)

Disparidad por diagnóstico: Llama3 (EN) (MH)




Disparidad por diagnóstico: Mistral (EN) (MH)






# Resultados del análisis de LLMs (02)

Disparidad por diagnóstico: Salamandra (EN) (MH)



Disparidad por diagnóstico: Salamandra (ES) (MH)



# Metodología de la mitigación de sesgo

Mitigación en MLMs (01)

Mitigación en LLMs (02)

# Mitigación en MLMs (01)



|                    |   |        |
|--------------------|---|--------|
| he<br>she          | → | they   |
| father<br>mother   | → | parent |
| Daniel<br>Isabella | → | [NAME] |
| 21m<br>27f         | → | age    |

*gender-scrubbing* sobre los datos de  
reentrenamiento

bert-base-uncased



SHBERT

bert-base-uncased



SHBERT-de-biased

Adaptación al dominio de  
autolesiones

# Mitigación en LLMs (02)




# Resultados de la mitigación de sesgo

Resultados de la mitigación en MLMs (01)

Resultados de la mitigación en LLMs (02)




# Resultados de la mitigación en MLMs (01)






# Resultados de la mitigación en LLMs (02)


Disparidad por diagnóstico: Llama3-3b-RLAIF-DPO (EN) (MH)



Disparidad por diagnóstico: Llama3-3b-SFT (EN) (MH)




# Conclusiones



anorexia  
depression

gambling addiction  
borderline personality disorder


Adaptación a dominio en MLMs



PTSD  
schizophrenia



cambio de idioma



TEPT  
esquizofrenia

Cambio de idioma en LLMs

SHBERT → SHBERT-de-biased

llama3.2:3b → llama-3-3b-SFT-gender-de-biased

*gender-scrubbing y SFT*

# Gracias

andreslilloortiz

